

Every AI Decision Is an Agency Decision

Why what AI takes on shapes what the organization becomes capable of

EMILY LIDDLE · CO-FOUNDER, THEORYA.AI

This Provocation applies Theory A, a management philosophy built on a single claim: human agency is the source of organizational value, and organizations either compound it or dissipate it. It is part of the Design for Agency Collection.

Work has always done two things. It produces the result the organization needs today, and it builds the people who will handle what comes next. AI is now taking on the work, and only the first is being counted.

Strip away the org chart, the systems, and the strategy decks, and what is left is an organization's ability to turn what it knows and what it can do into action. That is organizational agency. It draws on what its people see, judge, and decide, and on what its systems can now analyze and act on, and it depends on the organization being able to bring those together toward an outcome.

It shows up in specific, nameable capabilities: pricing a risk, reading a contract, seeing the opening a competitor left, knowing that the question everyone is working on is the wrong one. But organizational agency is more than a collection of abilities. It is what those capabilities become when the organization can connect them and put them to use toward an outcome, when the person who sees a problem can reach the person able to fix it, and when what one team learns changes how the next team works.

Organizations are deciding, right now, which of those capabilities keep being built. The decision is made in system settings, job designs, process changes, and choices about what AI takes on. **It shapes what the organization can do today and what it will remain able to do tomorrow.** That is a decision worth making deliberately.

01

Where Capability Comes From

Nobody starts out able to do the work an organization depends on. A new claims person works through hundreds of ordinary claims before they can spot the one that is wrong. A junior underwriter prices deal after deal before their judgment is worth

trusting. Doing the job is part of what teaches the job.

AI can accelerate that learning. It can explain, simulate, and put someone in front of far more examples than experience alone would have given them. Some of what used to take years genuinely does not need to anymore. But there is a difference between seeing a hundred difficult situations and being the person who had to decide one, be wrong, and live with what followed. **Exposure can be accelerated. Responsibility has to be carried.**

That responsibility develops capabilities that rarely appear in a training program or process map: sensing that something is off before being able to explain why, realizing that the problem everyone is solving is not the one worth solving, connecting two things because you happened to have carried both, trying something without a precedent because your judgment says it will work. These capabilities are formed through many conditions, but consequential work is one of the most powerful. It gives people a place to apply knowledge, make choices, receive feedback, and become answerable for what happens next.

The work was therefore doing two things. It produced the visible result: the claim paid, the deal priced, the question answered. It also contributed to building the person, the relationships, and the shared understanding the organization would need for the harder situation next year. The first shows up in every report. The second was rarely measured because the work kept producing it without anyone having to design for it.

Agency is not merely allocated through work. It is also produced through it.

02

The Ground Just Shifted

AI systems can now exercise broad agency inside knowledge work. They can notice a situation, interpret it, decide within defined bounds, and act. This is not the first form of automation, but it is a new scope and speed of artificial agency operating inside the organization.

That changes the design problem. An outcome tells the organization what the arrangement is meant to achieve. Human and artificial agency are the capacities it can bring to bear. How those forms of agency are orchestrated determines not only what gets done, but which capabilities grow, spread, narrow, or disappear. Agency compounds when the way work is arranged keeps making the organization better able to act.

Every decision about what work AI takes on is therefore two decisions at once. The first is about current return: what the organization can do now, how quickly, at what quality, and at what cost. The second is about what the organization will become capable of after the work has changed.

Both decisions are real, and they come from the same allocation. **Considered together, they are a design. Considered separately, the first one wins, because it is the one that appears in the business case.**

03

Why This Stays Invisible

The allocation rarely looks like a consequential design decision. It happens inside a process redesign, a job description, a budget, or a software feature turned on for reasons nobody remembers three years later.

Each local decision can be reasonable. Operations is reducing cycle time. Technology is adding functionality. Finance is protecting the return. A leader is rewriting a role around the work that remains. The question in front of each person is legitimate. What is usually absent is the question that crosses all of them: what will this arrangement leave the organization able to do?

The arrangement an organization inherits is often the sum of good answers to other questions.

Organizations do talk about capability, development, and succession. But that conversation usually happens somewhere else, with different people, after the work has already been divided. The allocation is made in operations, technology, and finance. Capability is considered later by people working with whatever that allocation left behind.

The money follows the same pattern. Every AI investment should be tested for return. The missing question is whether the same design is also building, preserving, or eliminating capabilities the organization will need after the payback period ends. Current return is evidence. Future capability begins as a hypothesis, so it is easier to leave outside the frame.

That is where the cost hides. An organization can hit every target it set while weakening its ability to respond to what comes next. The bill arrives later as work that must be rebuilt, expertise that must be bought back, decisions nobody can confidently make, or dependence on systems nobody can meaningfully challenge. **The organization does not see the loss while it is happening because every number it watched said yes the whole way down.**

04

The Line Keeps Moving

For much of the industrial era, organizations could treat the boundary between human and machine work as fixed between major redesigns. A machine changed what work people did, but its capability was stable enough to build jobs, processes, and careers around the division for years at a time.

AI changes that assumption. What it can take on moves with each model, each vendor release, and each discovery made inside the organization. The boundary between human agency, artificial agency, and shared agency does not stay settled. It moves continuously, and it is moving faster.

For an organization built around a one-time division of work, every movement becomes disruption. Roles have to be renegotiated, processes reopened, controls rewritten, and decisions everyone thought were finished made again. **The organization is redesigned on its vendors' schedule rather than through its own judgment.**

For an organization that knows where agency sits and why, the same movement becomes an advantage. When a capability changes, the question is narrower: what has changed, what does that make possible, and what should move now? The reasoning behind the current arrangement is visible. The people affected know how to participate. The organization can absorb the change and keep going.

The line will not stop moving. The design therefore cannot be static either.

05

Not Fewer People. Different Work.

This is not a promise that an organization will never need fewer people. **It is a refusal to make fewer people the design brief.**

Start with headcount and the question is how many people the current work needs. Start with the outcome and the question is what work should exist, who or what should do it, and what the arrangement should keep making possible. A smaller workforce may be one result. It cannot be the premise from which every other decision follows.

This is also not an argument to slow down on AI or protect work because people are attached to it. An organization that knows which capabilities it needs can hand AI more, not less. It can move with confidence because it knows what can transfer without consequence, what must remain available somewhere in the system, and what new work the released capacity makes possible.

Take a team whose job is answering questions and fixing problems for customers or employees. AI can handle most routine requests, and it should. The least-designed version leaves people with whatever remains: the escalations, exceptions, and hardest cases. The work becomes smaller, harder, and less sustainable by default.

A designed version starts with the outcome. If AI handles the routine, people can find and remove the causes that keep generating the same problems. They can shape where the service should go next. They can handle the moments where something real is at stake for the person on the other end and deepen the judgment those moments require. The workforce may change. But it changes because the work was

redesigned around the outcome, not because the old work was divided by a reduction target.

Because the boundary keeps moving, the answer cannot become a binder on a shelf. The organization has to be able to look again at how agency is divided, ask what the arrangement is producing and starving, and redesign while the choice is still changeable. Done once, this is a project. Done whenever the line moves, it becomes a capability the organization owns. Each redesign can leave it better able to make the next one.

06

The Quiet Clock

When a capability stops being built, nothing announces it. Experienced people may still be carrying the hard cases, and the results may hold. What disappears first is the pipeline behind them: fewer people developing the instinct that something is off, recognizing that the question is wrong, connecting what others have kept separate, or acting wisely without a precedent.

The organization discovers the loss when something unfamiliar arrives and nobody in the room has the judgment for it. That is the most expensive moment to begin rebuilding.

For now, there is another option, and it has a price. An organization can hire experience from people who developed theirs elsewhere, under older ways of working. But every company is drawing from the same pool while many of the conditions that filled it are disappearing. The stock of experienced people can remain visible even as the flow of new experts slows.

That is the clock. The loss is quiet, the early numbers can still look good, and the capability may eventually become unavailable at any price.

07

Deciding on Purpose

Design for Agency is a participatory method for deciding how human and artificial agency will be allocated and orchestrated around an outcome before that choice is buried in systems, roles, and budgets. It makes capability formation a criterion while the allocation can still be changed, so current return does not consume the organizational agency that future return depends on.

The starting point does not have to be the whole enterprise. It can be one outcome, one service, one operating boundary, or one live AI decision. The outcome sets the scope. The organization begins with what must become true, then identifies the work, decisions, capabilities, and forms of agency that contribute to it.

Some existing arrangements can also be reopened. If the allocation is still live or can still be changed, it remains a design question. Once the allocation cannot be changed and the capability loss has already occurred, the problem is no longer design. It is repair. Design for Agency can reveal that need, but the work then moves from allocation design to capability recovery: rebuilding apprenticeship, restoring judgment pathways, and recreating the conditions in which expertise can form.

The people currently doing the work must participate before the allocation is made, not receive a finished design afterward. They hold evidence no process map contains: where judgment actually enters, which exceptions matter, which ordinary experiences made them capable, and which workarounds keep the system functioning.

Their participation cannot become a way to extract that knowledge and make them easier to remove. People must control what they share. Participation and learning time must be funded. Released capacity must not automatically become a headcount reduction. Practitioner proposals need a real path to a decision that is explained. These are not courtesies around the method. They are conditions for getting honest evidence and a viable design.

Five questions reveal whether an arrangement was designed or simply arrived.

Why is the work set up this way? Someone should be able to explain the intent. When nobody can, the arrangement is usually the residue of decisions made for conditions that no longer exist.

Where does AI stop and human judgment begin? Not where the technology happens to stop today, but where the organization has decided the boundary should sit, and why.

Who remains answerable for the outcome? Accountability cannot disappear into the system. A named human authority or governing body must understand what the system is doing, remain answerable for the result, and be able to intervene.

What happens when conditions change, and what has authority to act? A person may step in. AI may adjust inside bounds someone has set. What matters is whether the response and the authority behind it were designed before the moment arrives.

Where is capability being built, and by what? People develop through experience, feedback, responsibility, teaching, and protected opportunities to stretch. Artificial capability improves through context, evaluation, updates, and governed feedback. The question is whether the whole arrangement is strengthening the capabilities the outcome will continue to require.

An intentional design has more than two choices. **Automate it or keep it is the poorest version of the decision.** For each capability, the organization can decide whether the system should compound it, create it, elevate it, distribute it, preserve it, transfer it, or retire it. Different capabilities within the same outcome may need different rulings. The ruling is not permanent. It is a hypothesis supported by current evidence, with conditions that trigger review as people, technology, and the environment change.

The test is double: does the design produce return now, and does it compound organizational agency over time? These are not competing scorecards. They are two requirements of the same design.

Design for Agency may begin as a deliberate engagement, but it works when it becomes how the organization revisits the division of work whenever the line moves. Continuous improvement asks how to make the work better. Design for Agency asks how human and artificial agency should work together now, what that arrangement is making the organization capable of, and when it needs to change again.

08

The Closing Provocation

Every organization is going to use AI, and a great deal of today's work should change or disappear. None of this is an argument for holding on to work that no longer creates value.

What is not settled is whether the organization will decide what the new arrangement should make possible and what capabilities it must continue to form. That decision is being made every day in process redesigns, job descriptions, system settings, budgets, and local choices made in good faith. The result will either be an arrangement the organization designed or one it inherited.

The gap between those organizations is small at first. One quarter apart, they may look identical. The difference appears when the line moves again. One can absorb the next shift, reconsider the allocation, and redirect agency toward the outcome. The other has to discover what it lost before it can begin to respond. That gap compounds every time.

Every decision about what work AI takes on is also a decision about whether the organization pulls ahead or quietly falls behind. There is no version in which the decision does not get made. There is only the version in which the organization makes it deliberately, and the version in which it inherits the answer.

This Provocation introduces the Design for Agency Collection, one of the Expertise Collections of TheoryA.ai.